Package 'car2'

July 13, 2015

Type Package

Title Extends capabilities of 'car' to include companion functions for logistic regression

Version 0.1

Date 2015-05-26

Author Selva Prabhakaran

Maintainer Selva Prabhakaran <selva86@gmail.com>

Description Provides companion function for analysing the performance of classification models, based on the problems specific objectives. There is a function to plot the ROC curve on the beautiful ggplot2 graphics framework, computing AUROC, concordance, discordance, specificity, sensitivity, Youden's index, Somers D statistic etc.

License GPL (>= 2)

LazyData TRUE

LazyLoad yes

Depends ggplot2

Import ggplot2

R topics documented:

ActualsAndScores *ActualsAndScores*

Description

A dataset containing the actuals for a simulated binary response variable as a numeric and the prediction probablity scores for a classification model like logistic regression.

Usage

```
data(ActualsAndScores)
```

Format

A data frame with 170 rows and 2 variables

Details

- Actuals. A simulated variable meant to serve as the actual binary response variable. The good/events are marked as 1 while the bads/non-events are marked 0.
- PredictedScores. The prediction probability scores based on a classification model.

AUROC *AUROC*

Description

Calculate the area uder ROC curve statistic for a given logit model.

Usage

```
AUROC(actuals, predictedScores)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.

Details

For a given actuals and predicted probability scores, the area under the ROC curve shows how well the model performs at capturing the false events and false non-events. An best case model will have an area of 1. However that would be unrealistic, so the closer the aROC to 1, the better is the model.

Value

The area under the ROC curve for a given logit model.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
AUROC(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

Concordance *Concordance*

Description

Calculate concordance and discordance percentages for a logit model

Usage

```
Concordance(actuals, predictedScores)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.

Details

Calculate the percentage of concordant and discordant pairs for a given logit model.

Value

a list containing percentage of concordant pairs, percentage discordant pairs, percentage ties and No. of pairs.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
Concordance(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

confusionMatrix *confusionMatrix*

Description

Calculate the confusion matrix for the fitted values for a logistic regression model.

Usage

```
confusionMatrix(actuals, predictedScores, threshold = 0.5)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.
threshold	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.

Details

For a given actuals and predicted probability scores, the confusion matrix showing the count of predicted events and non-events against actual events and non events.

Value

For a given actuals and predicted probability scores, returns the confusion matrix showing the count of predicted events and non-events against actual events and non events.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
confusionMatrix(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

kappaCohen *kappaCohen*

Description

Calculate the Cohen's kappa statistic for a given logit model.

Usage

```
kappaCohen(actuals, predictedScores, threshold = 0.5)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.
threshold	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.

Details

For a given actuals and predicted probability scores, Cohen's kappa is calculated. Cohen's kappa is calculated as (probabiliity of agreement - probability of expected) / (1-(probability of expected)))

Value

The Cohen's kappa of the given actuals and predicted probability scores

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
kappaCohen(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

misClassError *misClassError*

Description

Calculate the percentage misclassification error for this logit model's fitted values.

Usage

```
misClassError(actuals, predictedScores, threshold = 0.5)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.
threshold	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.

Details

For a given binary response actuals and predicted probability scores, misclassfication error is the number of mismatches between the predicted and actuals direction of the binary y variable.

Value

The misclassification error, which tells what proportion of predicted direction did not match with the actuals.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
misClassError(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores, threshold
```

---

plotROC *plotROC*

Description

Plot the Receiver Operating Characteristics(ROC) Curve based on ggplot2

Usage

```
plotROC(actuals, predictedScores, threshold = 0.5, Show.labels = F)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.
threshold	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.
Show.labels	Whether the probability scores should be printed at change points?. Defaults to False.

Details

For a given actuals and predicted probability scores, A ROC curve is plotted using the ggplot2 framework along the the area under the curve.

Value

Plots the ROC curve

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
plotROC(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

`sensitivity` *sensitivity*

Description

Calculate the sensitivity for a given logit model.

Usage

```
sensitivity(actuals, predictedScores, threshold = 0.5)
```

Arguments

`actuals`	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
`predictedScores`	The prediction probability scores for each observation.
`threshold`	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.

Details

For a given binary response actuals and predicted probability scores, sensitivity is defined as number of observations with the event AND predicted to have the event divided by the number of observations with the event. It can be used as an indicator to gauge how sensitive is your model in detecting the occurence of events, especially when you are not so concerned about predicting the non-events as true.

Value

The sensitivity of the given binary response actuals and predicted probability scores, which is, the number of observations with the event AND predicted to have the event divided by the nummber of observations with the event.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
sensitivity(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

---

somersD *somersD*

Description

Calculate the Somers D statistic for a given logit model

Usage

```
somersD(actuals, predictedScores)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.

Details

For a given binary response actuals and predicted probability scores, Somer's D is calculated as the number of concordant pairs less number of discordant pairs divided by total number of pairs.

Value

The Somers D statistic, which tells how many more concordant than discordant pairs exist divided by total number of pairs.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
somersD(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

specificity *specificity*

Description

Calculate the specificity for a given logit model.

Usage

```
specificity(actuals, predictedScores, threshold = 0.5)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.
threshold	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.

Details

For a given given binary response actuals and predicted probability scores, specificity is defined as number of observations without the event AND predicted to not have the event divided by the number of observations without the event. Specificity is particularly useful when you are extra careful not to predict a non event as an event, like in spam detection where you dont want to classify a genuine mail as spam(event) where it may be somewhat ok to occasionally classify a spam as a genuine mail(a non-event).

Value

The specificity of the given binary response actuals and predicted probability scores, which is, the number of observations without the event AND predicted to not have the event divided by the nummber of observations without the event.

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
specificity(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

youdensIndex *youdensIndex*

Description

Calculate the specificity for a given logit model.

Usage

```
youdensIndex(actuals, predictedScores, threshold = 0.5)
```

Arguments

actuals	The actual binary flags for the response variable. It can take values of either 1 or 0, where 1 represents the 'Good' or 'Events' while 0 represents 'Bad' or 'Non-Events'.
predictedScores	The prediction probability scores for each observation.
threshold	If predicted value is above the threshold, it will be considered as an event (1), else it will be a non-event (0). Defaults to 0.5.

Details

For a given binary response actuals and predicted probability scores, Youden's index is calculated as sensitivity + specificity - 1

Value

The youdensIndex of the given binary response actuals and predicted probability scores, which is calculated as Sensitivity + Specificity - 1

Author(s)

Selva Prabhakaran

Examples

```
data('ActualsAndScores')
youdensIndex(actuals=ActualsAndScores$Actuals, predictedScores=ActualsAndScores$PredictedScores)
```

Index